

Genotype quality control with plinkQC

Hannah Meyer

2025-11-07

Contents

Introduction	1
Per-individual quality control	2
Per-marker quality control	2
Clean data	2
Workflow	2
Per-individual quality control	3
Per-marker quality control	6
Create QC-ed dataset	9
Step-by-step	9
Individuals with discordant sex information	9
Individuals with outlying missing genotype and/or heterozygosity rates	10
Related individuals	11
Individuals of divergent ancestry	12
Markers with excessive missingness rate	13
Markers with deviation from HWE	14
Markers with low minor allele frequency	15
References	16

Introduction

Genotyping arrays enable the direct measurement of an individual's genotype at thousands of markers. Subsequent analyses such as genome-wide association studies rely on the high quality of these marker genotypes.

Anderson and colleagues introduced a protocol for data quality control in genetic association studies heavily based on the summary statistics and relatedness estimation functions in the PLINK software [1]. PLINK is a comprehensive, open-source command-line tool for genome-wide association studies (GWAS) and population genetics research [3]. Its main functionalities include data management, computing individual- and marker-level summary statistics, identity-by-state estimation and association analysis.

Integration with R is achieved through its R plugin or PLINK/SEQ R Package [4]. While the plugin is limited to operations yielding simple genetic marker vectors as output, the PLINK/SEQ R Package is limited in the functionalities it can access.

plinkQC facilitates genotype quality control for genetic association studies as described by [1]. It wraps around plink basic statistics (e.g. missing genotyping rates per individual, allele frequencies per genetic marker) and relationship functions and generates a per-individual and per-marker quality control report. Individuals and markers that fail the quality control can subsequently be removed with *plinkQC* to generate a new, clean dataset. Removal of individuals based on relationship status is optimised to retain as many individuals as possible in the study.

plinkQC depends on the PLINK (**version 1.9**), which has to be manually installed prior to the usage of *plinkQC*. It assumes the genotype have already been determined from the original probe intensity data of the genotype array and is available in plink format.

Note: PLINK 2.0 is still in alpha status with many re-implementations and updates, including output file changes. While these changes are ongoing, *plinkQC* will rely on users using PLINK 1.9. I will monitor PLINK 2.0 changes and check compatibility with *plinkQC*. If users require specific PLINK 2.0 output, I recommend using the Step-by-Step approach described below and manually saving to the output files expected from PLINK 1.9.

The protocol is implemented in three main functions, the per-individual quality control (`perIndividualQC`), the per-marker quality control (`perMarkerQC`) and the generation of the new, quality control dataset (`cleanData`):

Per-individual quality control

The per-individual quality control with `perIndividualQC` wraps around these functions: (i) `check_sex`: for the identification of individuals with discordant sex information, (ii) `check_heterozygosity_and_missingness`: for the identification of individuals with outlying missing genotype and/or heterozygosity rates, (iii) `check_relatedness`: for the identification of related individuals, (iv) `check_ancestry`: identification of individuals of divergent ancestry.

Per-marker quality control

The per-marker quality control with `perMarkerQC` wraps around these functions: (i) `check_snp_missingnes`: for the identifying markers with excessive missing genotype rates, (ii) `check_hwe`: for the identifying markers showing a significant deviation from Hardy-Weinberg equilibrium (HWE), (iii) `check_maf`: for the removal of markers with low minor allele frequency (MAF).

Clean data

`cleanData` takes the results of `perMarkerQC` and `perIndividualQC` and creates a new dataset with all individuals and markers that passed the quality control checks.

Workflow

In the following, genotype quality control with *plinkQC* is applied on a small example dataset with 200 individuals and 10,000 markers (provided with this package). The quality control is demonstrated in three easy steps, per-individual and per-marker quality control followed by the generation of the new dataset. In addition, the functionality of each of the functions underlying `perMarkerQC` and `perIndividualQC` is demonstrated at the end of this vignette.

```

package.dir <- find.package('plinkQC')
indir <- file.path(package.dir, 'extdata')
qcdir <- tempdir()
name <- 'data'
path2plink <- "/Users/hannah/bin/plink"

```

Per-individual quality control

For `perIndividualQC`, one simply specifies the directory where the data is stored (`qcdir`) and the prefix of the plink files (i.e. `prefix.bim`, `prefix.bed`, `prefix.fam`). In addition, the names of the files containing information about the reference population and the merged dataset used in `check_ancestry` have to be provided: `refSamplesFile`, `refColorsFile` and `prefixMergedDataset`. Per default, all quality control checks will be conducted.

In addition to running each check, `perIndividualQC` writes a list of all fail individual IDs to the `qcdir`. These IDs will be removed in the computation of the `perMarkerQC`. If the list is not present, `perMarkerQC` will send a message about conducting the quality control on the entire dataset.

NB: To reduce the data size of the example data in `plinkQC`, `data.genome` has already been reduced to the individuals that are related. Thus the relatedness plots in C only show counts for related individuals only.

NB: To demonstrate the results of the ancestry check, the required eigenvector file of the combined study and reference datasets have been precomputed and for the purpose of this example will be copied to the `qcdir`. In practice, the `qcdir` will often be the same as the `indir` and this step will not be required.

```

file.copy(file.path(package.dir, 'extdata/data.HapMapIII.eigenvec'), qcdir)
#> [1] TRUE

```

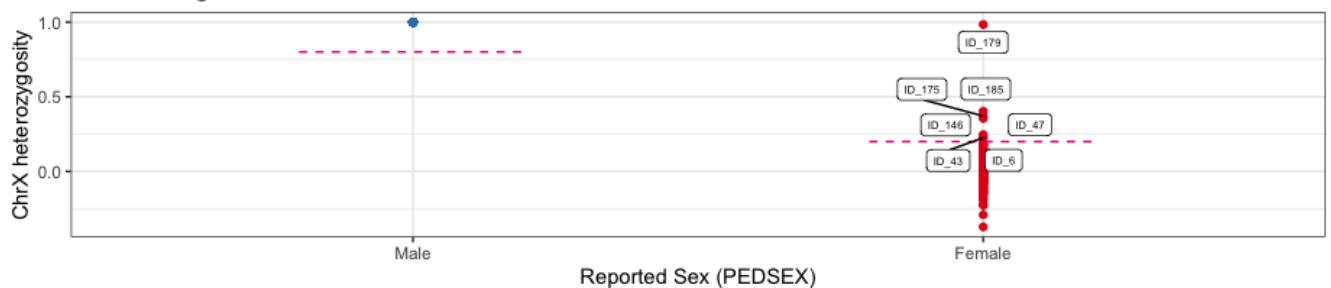
`perIndividualQC` displays the results of the quality control steps in a multi-panel plot.

```

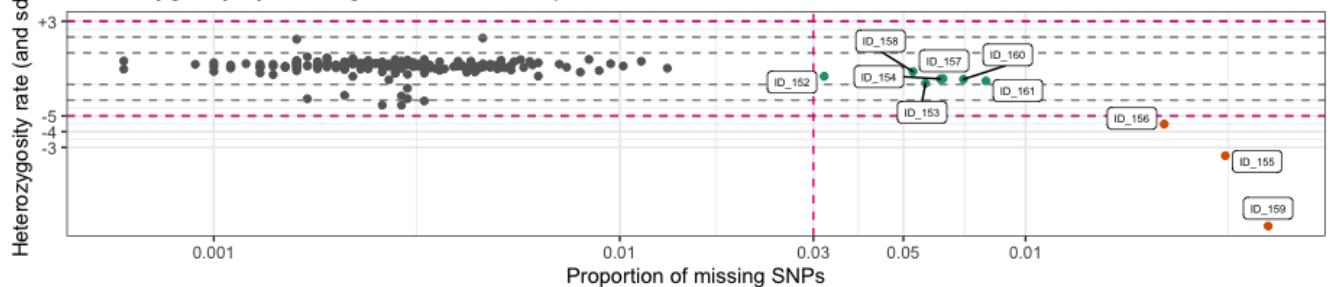
fail_individuals <- perIndividualQC(indir=indir, qcdir=qcdir, name=name,
                                      refSamplesFile=paste(indir, "/HapMap_ID2Pop.txt",
                                                           sep=""),
                                      refColorsFile=paste(indir, "/HapMap_PopColors.txt",
                                                           sep=""),
                                      prefixMergedDataset="data.HapMapIII",
                                      path2plink=path2plink, do.run_check_ancestry = FALSE,
                                      interactive=TRUE, verbose=TRUE)

```

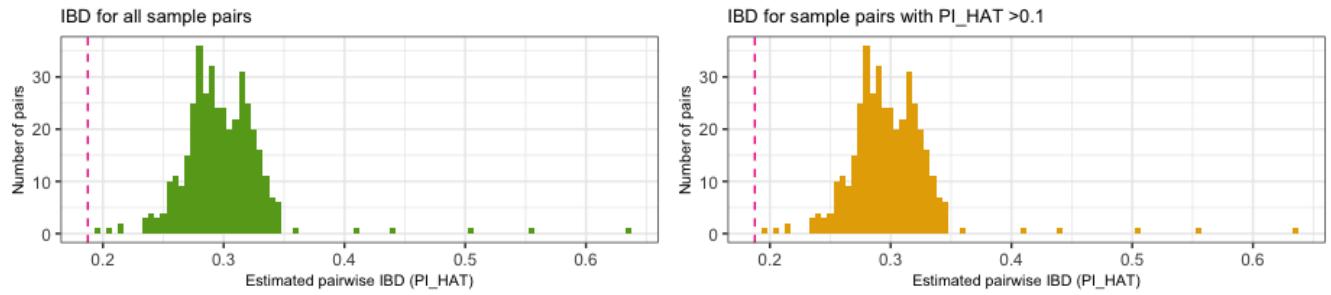
A Check assigned sex versus SNP sex



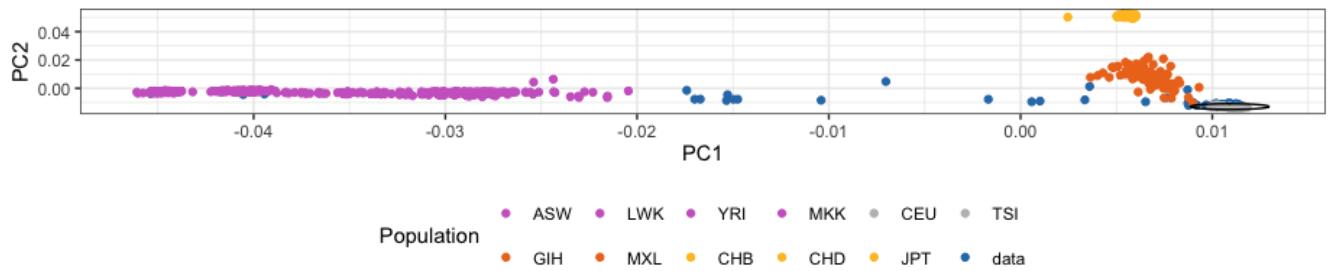
B Heterozygosity by Missingness across samples



C Relatedness estimated as pairwise IBD (PI_HAT)

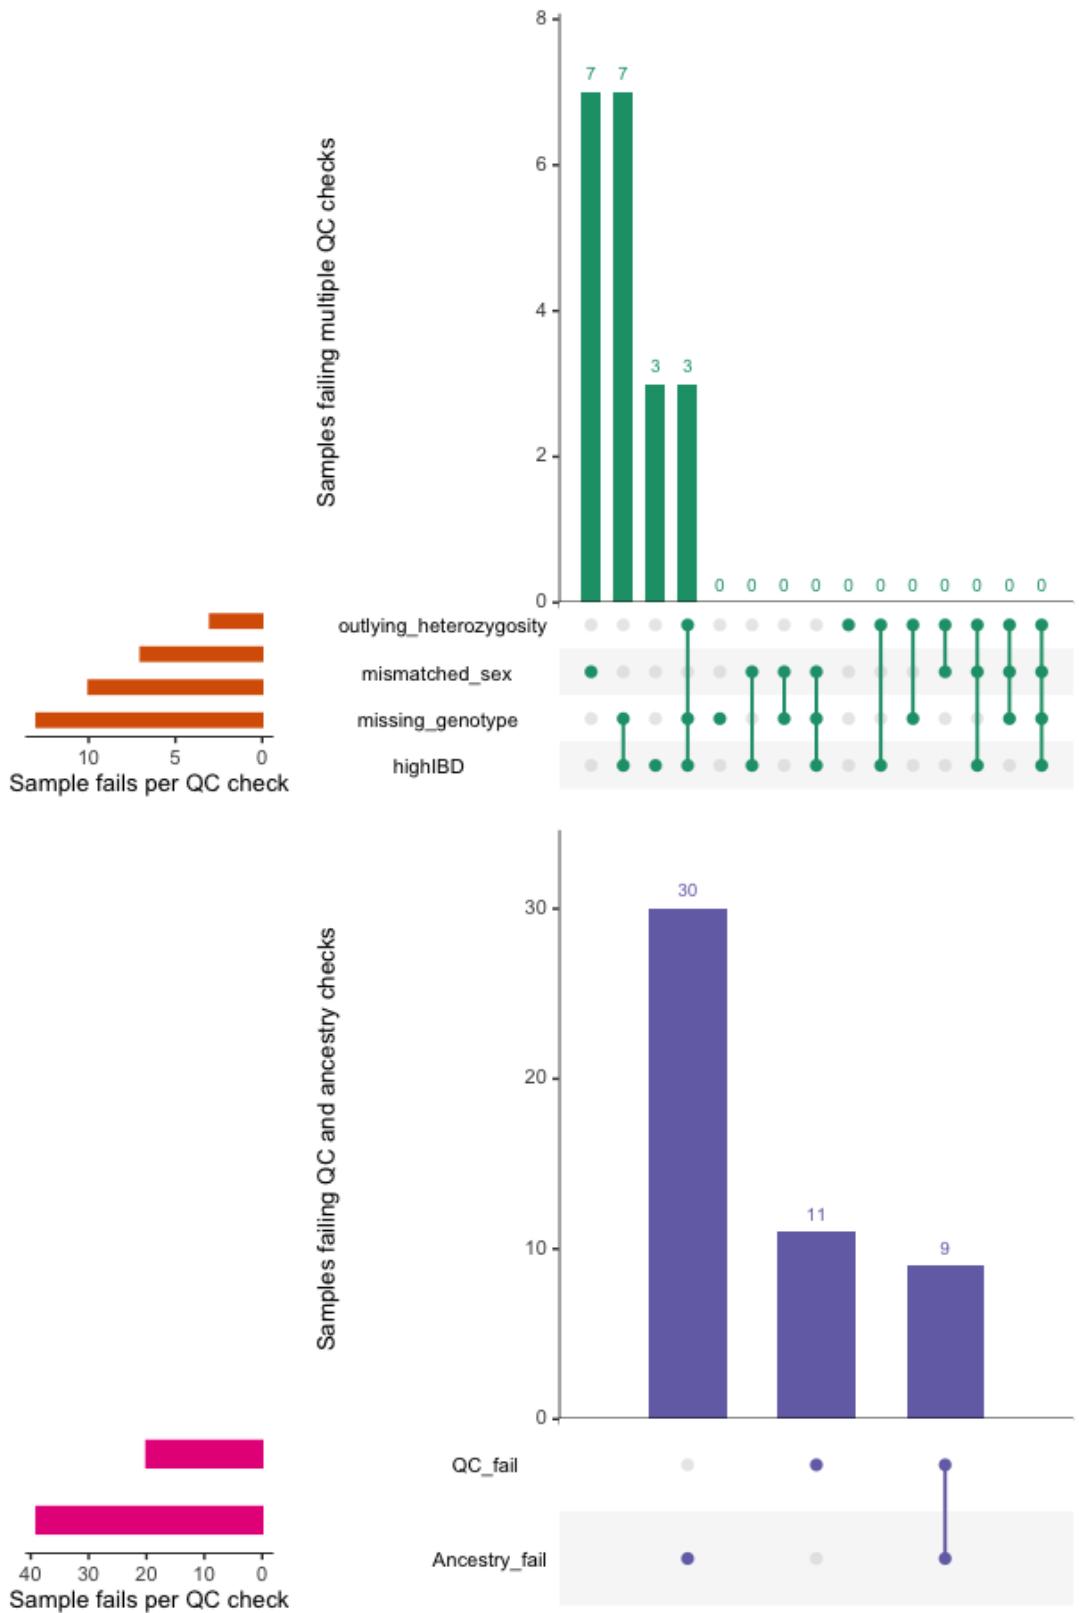


D PCA on combined reference and study genotypes



overviewperIndividualQC depicts overview plots of quality control failures and the intersection of quality control failures with ancestry exclusion.

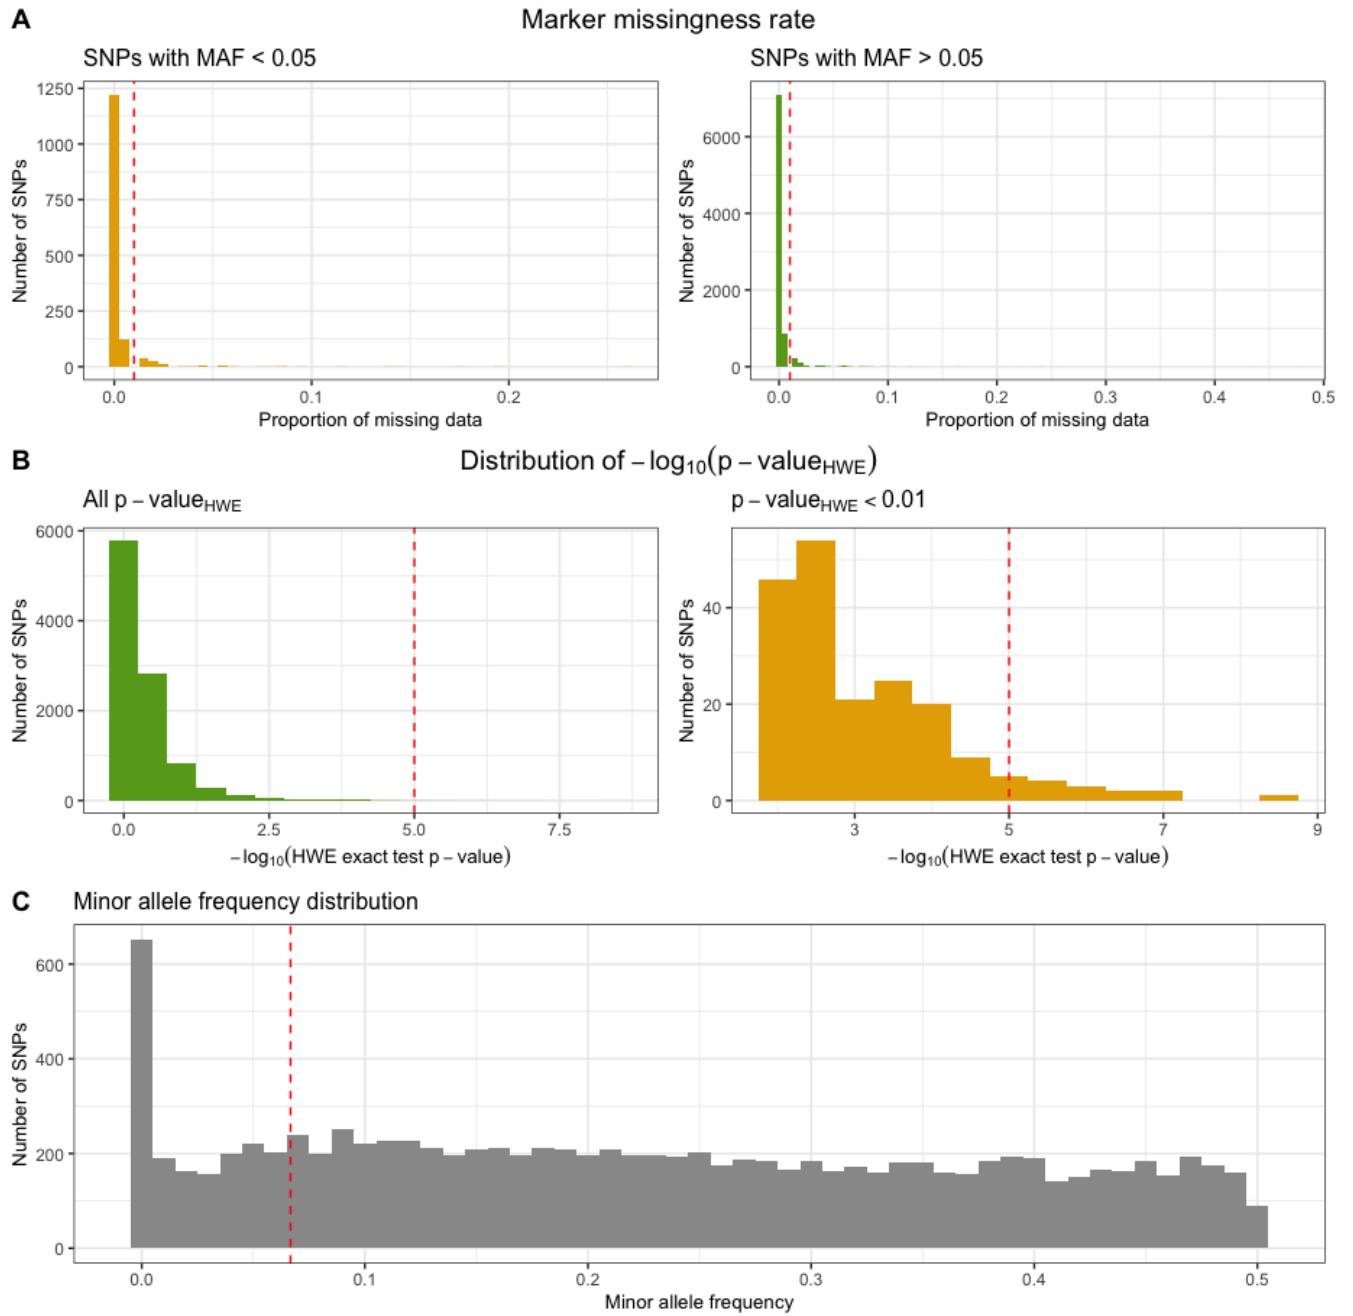
```
overview_individuals <- overviewPerIndividualQC(fail_individuals,
                                                interactive=TRUE)
```



Per-marker quality control

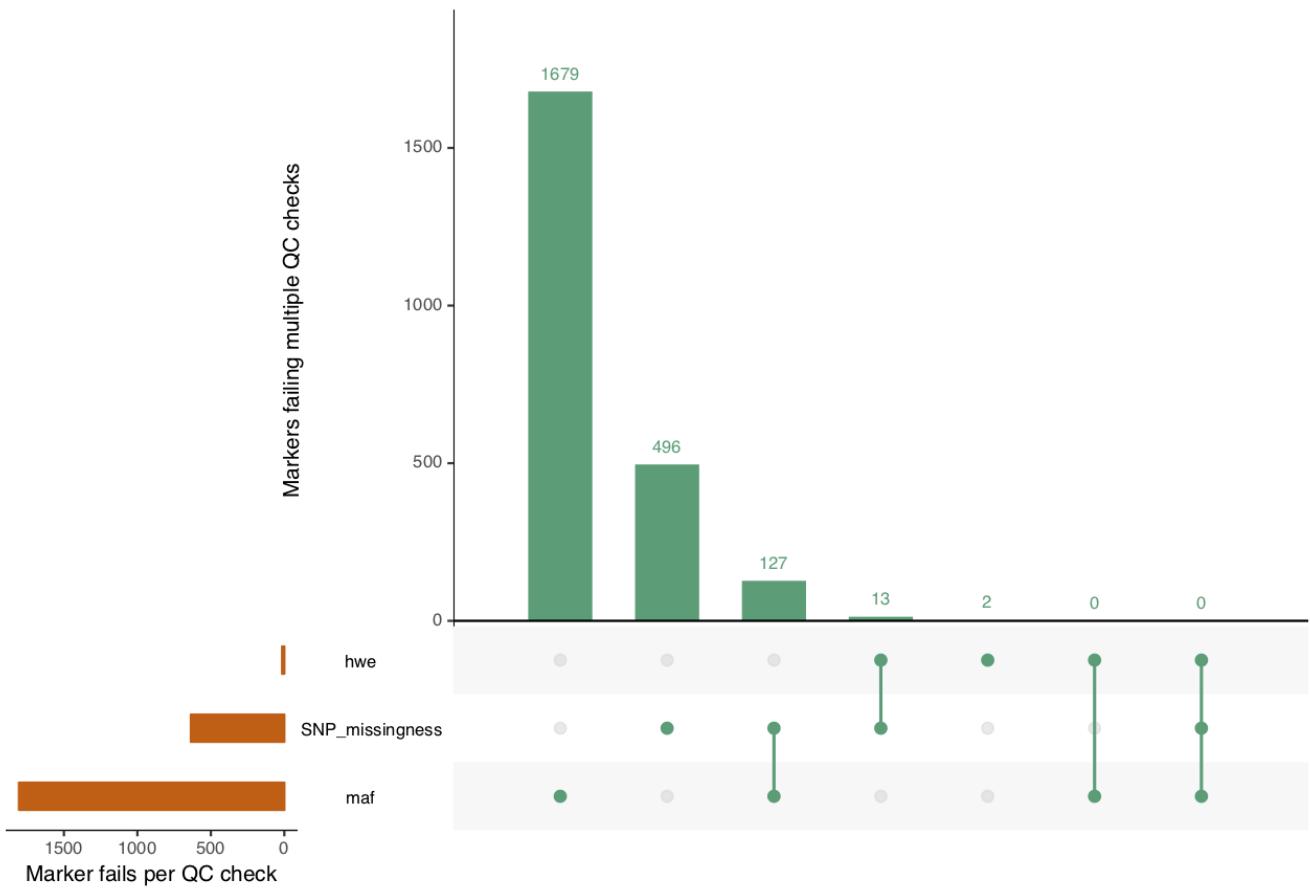
`perMarkerQC` applies its checks to data in the specified directory (`qmdir`), starting with the specified prefix of the plink files (i.e. `prefix.bim`, `prefix.bed`, `prefix.fam`). Optionally, the user can specify different thresholds for the quality control checks and which check to conduct. Per default, all quality control checks will be conducted. `perMarkerQC` displays the results of the QC step in a multi-panel plot.

```
fail_markers <- perMarkerQC(indir=indir, qmdir=qmdir, name=name,
                             path2plink=path2plink,
                             verbose=TRUE, interactive=TRUE,
                             showPlinkOutput=FALSE)
```



`overviewPerMarkerQC` depicts an overview of the marker quality control failures and their overlaps.

```
overview_marker <- overviewPerMarkerQC(fail_markers, interactive=TRUE)
```



Create QC-ed dataset

After checking results of the per-individual and per-marker quality control, individuals and markers that fail the chosen criteria can automatically be removed from the dataset with `cleanData`, resulting in the new dataset `qmdir/data.clean.bed`, `qmdir/data.clean.bim`, `qmdir/data.clean.fam`. For convenience, `cleanData` returns a list of all individuals in the study split into keep and remove individuals.

```
Ids <- cleanData(indir=indir, qmdir=qmdir, name=name, path2plink=path2plink,  
    verbose=TRUE, showPlinkOutput=FALSE)
```

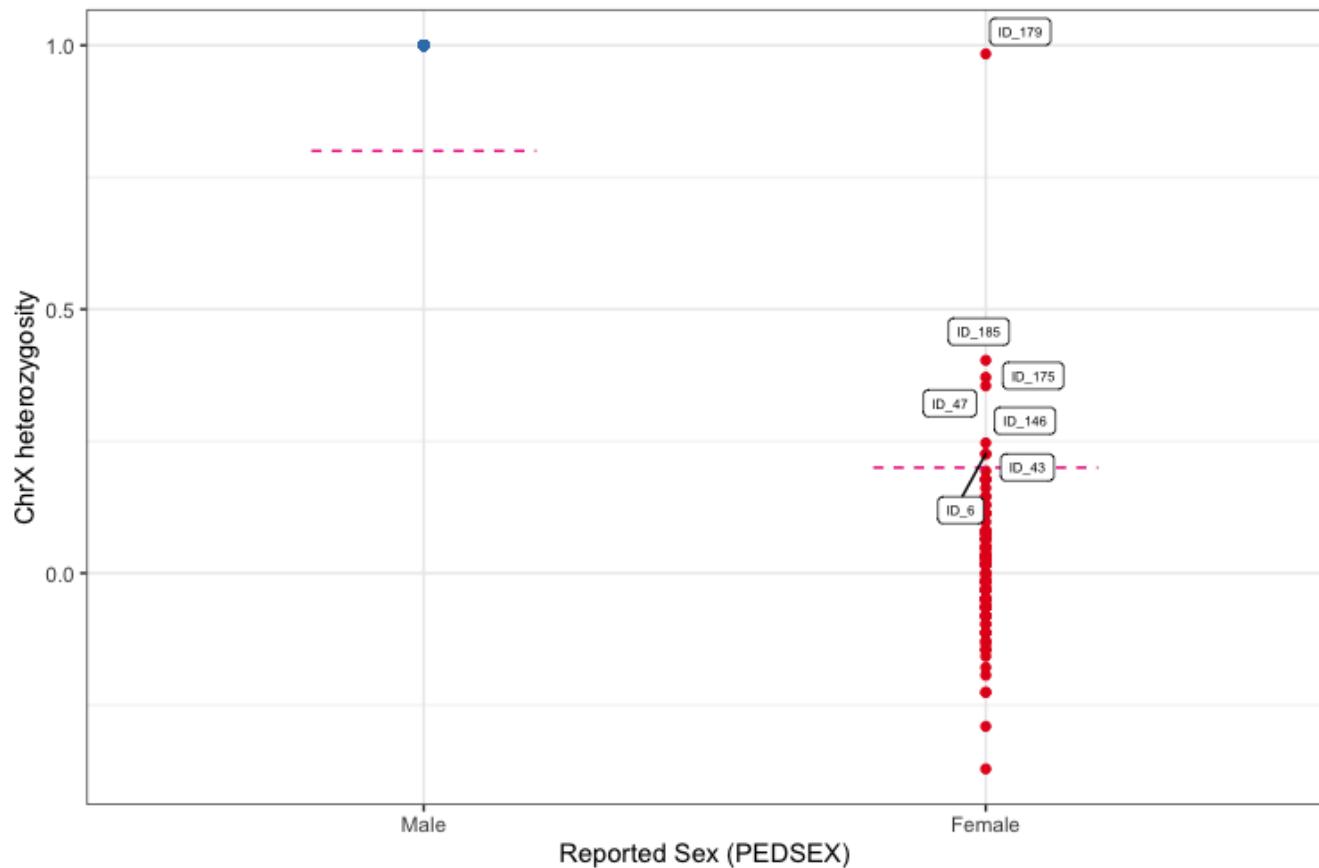
Step-by-step

Individuals with discordant sex information

The identification of individuals with discordant sex information helps to detect sample mix-ups and samples with very poor genotyping rates. For each sample, the homozygosity rates across all X-chromosomal genetic markers are computed and compared with the expected rates (typically $\$<\0.2 for females and $\$>\0.8 for males). For samples where the assigned sex (`PEDSEX` in the `.fam` file) contradicts the sex inferred from the homozygosity rates (`SNPSEX`), it should be checked that the sex was correctly recorded (genotyping often occurs at different locations as phenotyping and misrecording might occur). Samples with discordant sex information that is not accounted for should be removed from the study. Identifying individuals with discordant sex information is implemented in `check_sex`. It finds individuals whose `SNPSEX` \neq `PEDSEX`. Optionally, an extra `data.frame` with sample IDs and sex can be provided to double check if external and `PEDSEX` data (often processed at different centers) match. If a mismatch between `PEDSEX` and `SNPSEX` was detected, by `SNPSEX == Sex`, `PEDSEX` of these individuals can optionally be updated. `check_sex` depicts the X-chromosomal heterozygosity (`SNPSEX`) of the samples split by their (`PEDSEX`).

```
fail_sex <- check_sex(indir=indir, qmdir=qmdir, name=name, interactive=TRUE,  
    verbose=TRUE, path2plink=path2plink)
```

Check assigned sex versus SNP sex

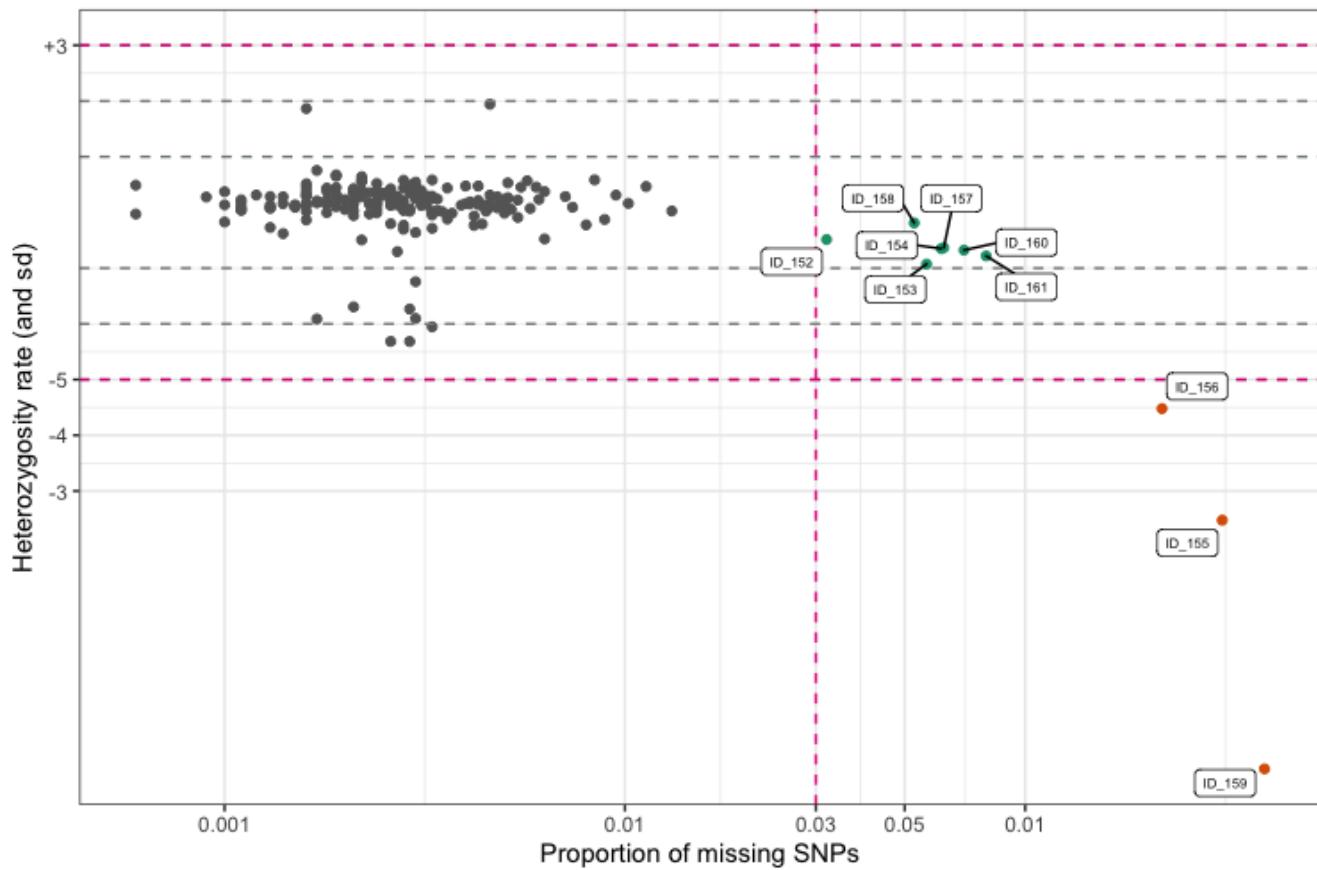


Individuals with outlying missing genotype and/or heterozygosity rates

The identification of individuals with outlying missing genotype and/or heterozygosity rates helps to detect samples with poor DNA quality and/or concentration that should be excluded from the study. Typically, individuals with more than 3-7% of their genotype calls missing are removed. Outlying heterozygosity rates are judged relative to the overall heterozygosity rates in the study, and individuals whose rates are more than a few standard deviations (sd) from the mean heterozygosity rate are removed. A typical quality control for outlying heterozygosity rates would remove individuals who are three sd away from the mean rate. Identifying related individuals with outlying missing genotype and/or heterozygosity rates is implemented in `check_het_and_miss`. It finds individuals that have genotyping and heterozygosity rates that fail the set thresholds and depicts the results as a scatter plot with the samples' missingness rates on x-axis and their heterozygosity rates on the y-axis.

```
fail_het_imiss <- check_het_and_miss(indir=indir, qmdir=qmdir, name=name,
                                         interactive=TRUE, path2plink=path2plink)
```

Heterozygosity by Missingness across samples



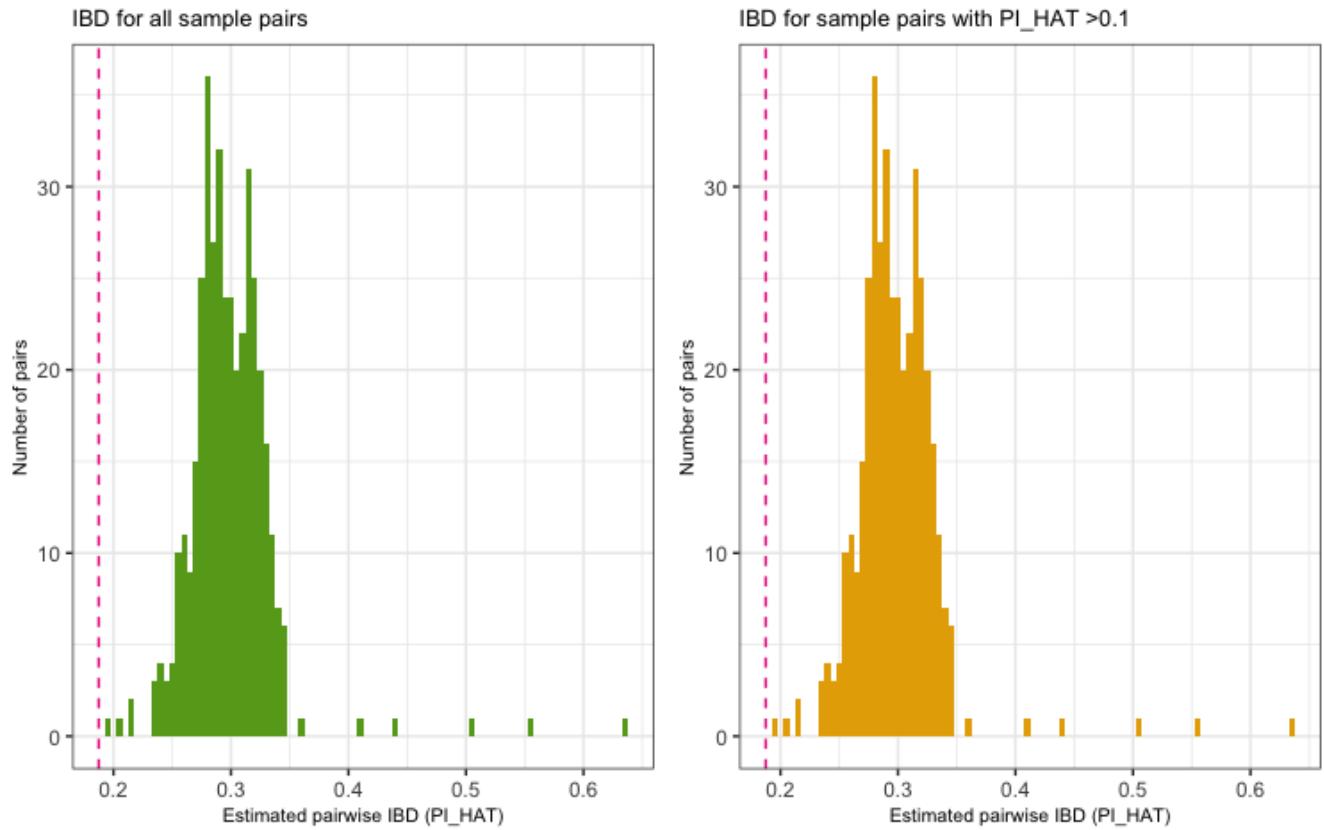
Related individuals

Depending on the future use of the genotypes, it might required to remove any related individuals from the study. Related individuals can be identified by their proportion of shared alleles at the genotyped markers (identity by descend, IBD). Standardly, individuals with second-degree relatedness or higher will be excluded. Identifying related individuals is implemented in `check_relatedness`. It finds pairs of samples whose proportion of IBD is larger than the specified `highIBDTh`. Subsequently, for pairs of individual that do not have additional relatives in the dataset, the individual with the greater genotype missingness rate is selected and returned as the individual failing the relatedness check. For more complex family structures, the unrelated individuals per family are selected (e.g. in a parents-offspring trio, the offspring will be marked as fail, while the parents will be kept in the analysis).

NB: To reduce the data size of the example data in `plinkQC`, `data.genome` has already been reduced to the individuals that are related. Thus the relatedness plots in C only show counts for related individuals only.

```
exclude_relatedness <- check_relatedness(indir=indir, qcdir=qcdir, name=name,
                                         interactive=TRUE,
                                         path2plink=path2plink)
```

Relatedness estimated as pairwise IBD (PI_HAT)

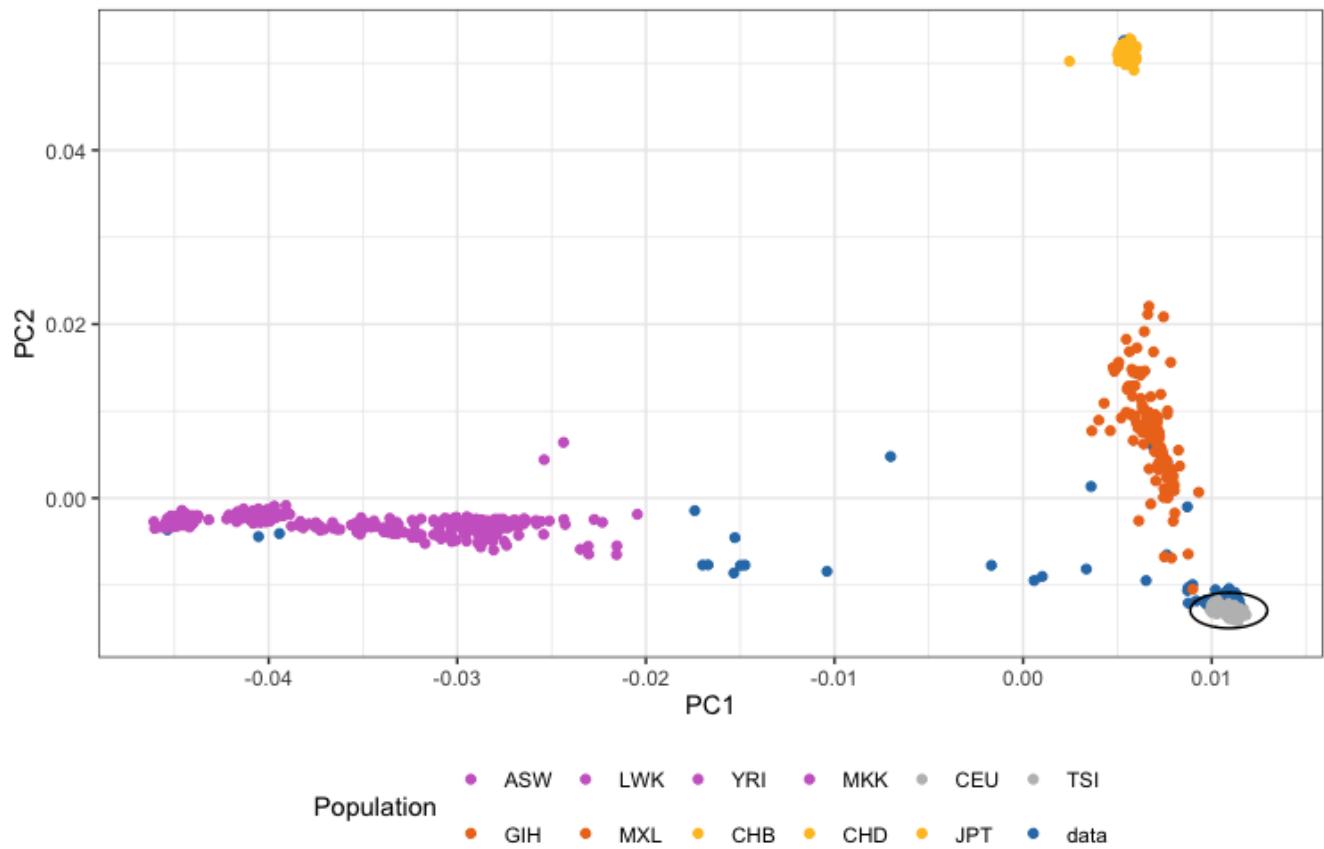


Individuals of divergent ancestry

The identification of individuals of divergent ancestry can be achieved by combining the genotypes of the study population with genotypes of a reference dataset consisting of individuals from known ethnicities (for instance individuals from the Hapmap or 1000 genomes study [9]). Principal component analysis on this combined genotype panel can be used to detect population structure down to the level of the reference dataset (for Hapmap and 1000 Genomes, this is down to large-scale continental ancestry). Identifying individuals of divergent ancestry is implemented in `check_ancestry`. Currently, check ancestry only supports automatic selection of individuals of European descent. It uses information from principal components 1 and 2 to find the center of the European reference samples. All study samples whose euclidean distance from the centre falls outside a specified radius are considered non-European. `check_ancestry` creates a scatter plot of PC1 versus PC2 color-coded for samples of the reference populations and the study population.

```
exclude_ancestry <- check_ancestry(indir=indir, qmdir=qmdir, name=name,
                                    refSamplesFile=paste(indir, "/HapMap_ID2Pop.txt",
                                                         sep=""),
                                    refColorsFile=paste(indir, "/HapMap_PopColors.txt",
                                                         sep=""),
                                    prefixMergedDataset="data.HapMapIII",
                                    path2plink=path2plink, run.check_ancestry = FALSE,
                                    interactive=TRUE)
```

PCA on combined reference and study genotypes

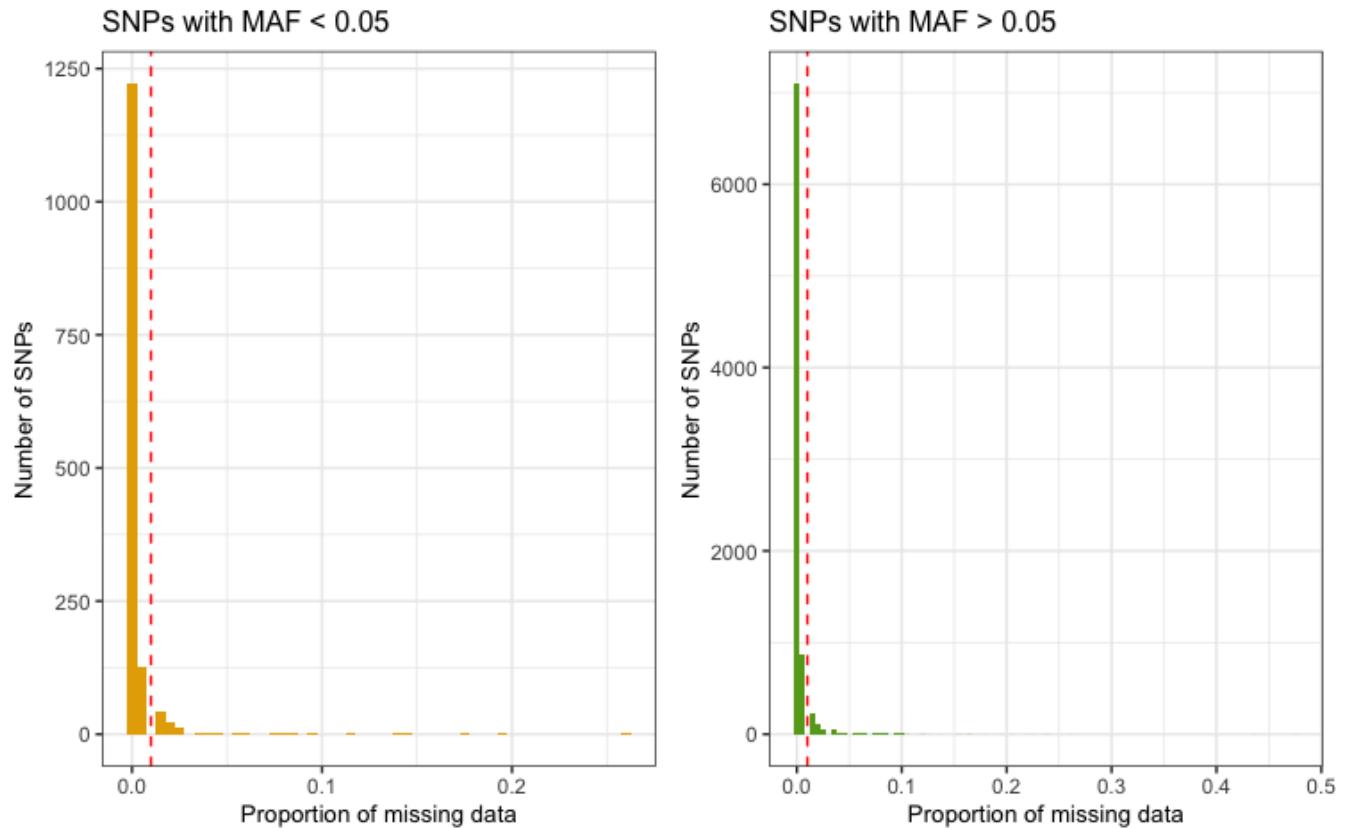


Markers with excessive missingness rate

Markers with excessive missingness rate are removed as they are considered unreliable. Typically, thresholds for marker exclusion based on missingness range from 1%-5%. Identifying markers with high missingness rates is implemented in `snp_missingness`. It calculates the rates of missing genotype calls and frequency for all variants in the individuals that passed the `perIndividualQC`.

```
fail_snpmissing <- check.snp_missingness(indir=indir, qmdir=qmdir, name=name,
                                         interactive=TRUE,
                                         path2plink=path2plink,
                                         showPlinkOutput=FALSE)
```

Marker missingness rate

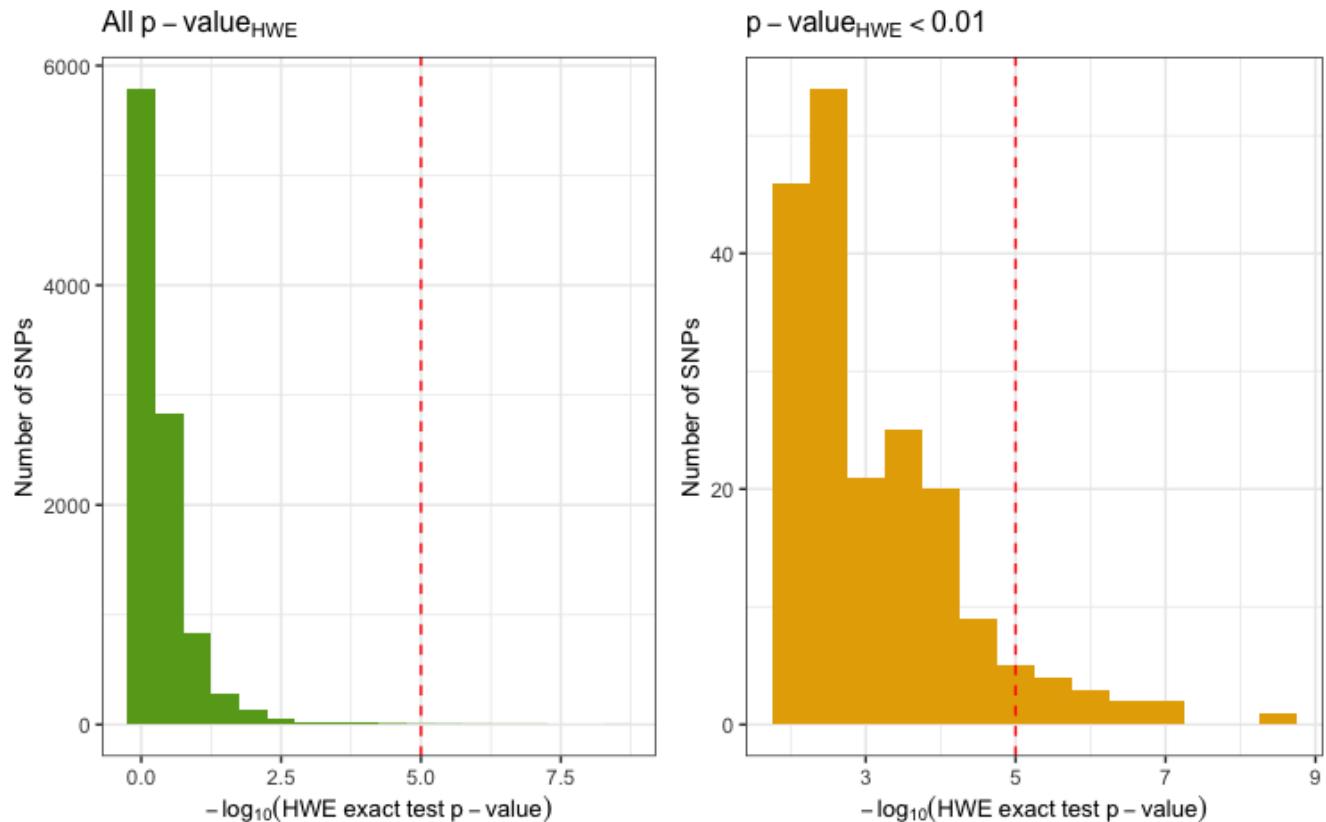


Markers with deviation from HWE

Markers with strong deviation from HWE might be indicative of genotyping or genotype-calling errors. As serious genotyping errors often yield very low p-values (in the order of 10^{-50}), it is recommended to choose a reasonably low threshold to avoid filtering too many variants (that might have slight, non-critical deviations). Identifying markers with deviation from HWE is implemented in `check_hwe`. It calculates the observed and expected heterozygote frequencies per SNP in the individuals that passed the `perIndividualQC` and computes the deviation of the frequencies from Hardy-Weinberg equilibrium (HWE) by HWE exact test.

```
fail_hwe <- check_hwe(indir=indir, qmdir=qmdir, name=name, interactive=TRUE,  
path2plink=path2plink, showPlinkOutput=FALSE)
```

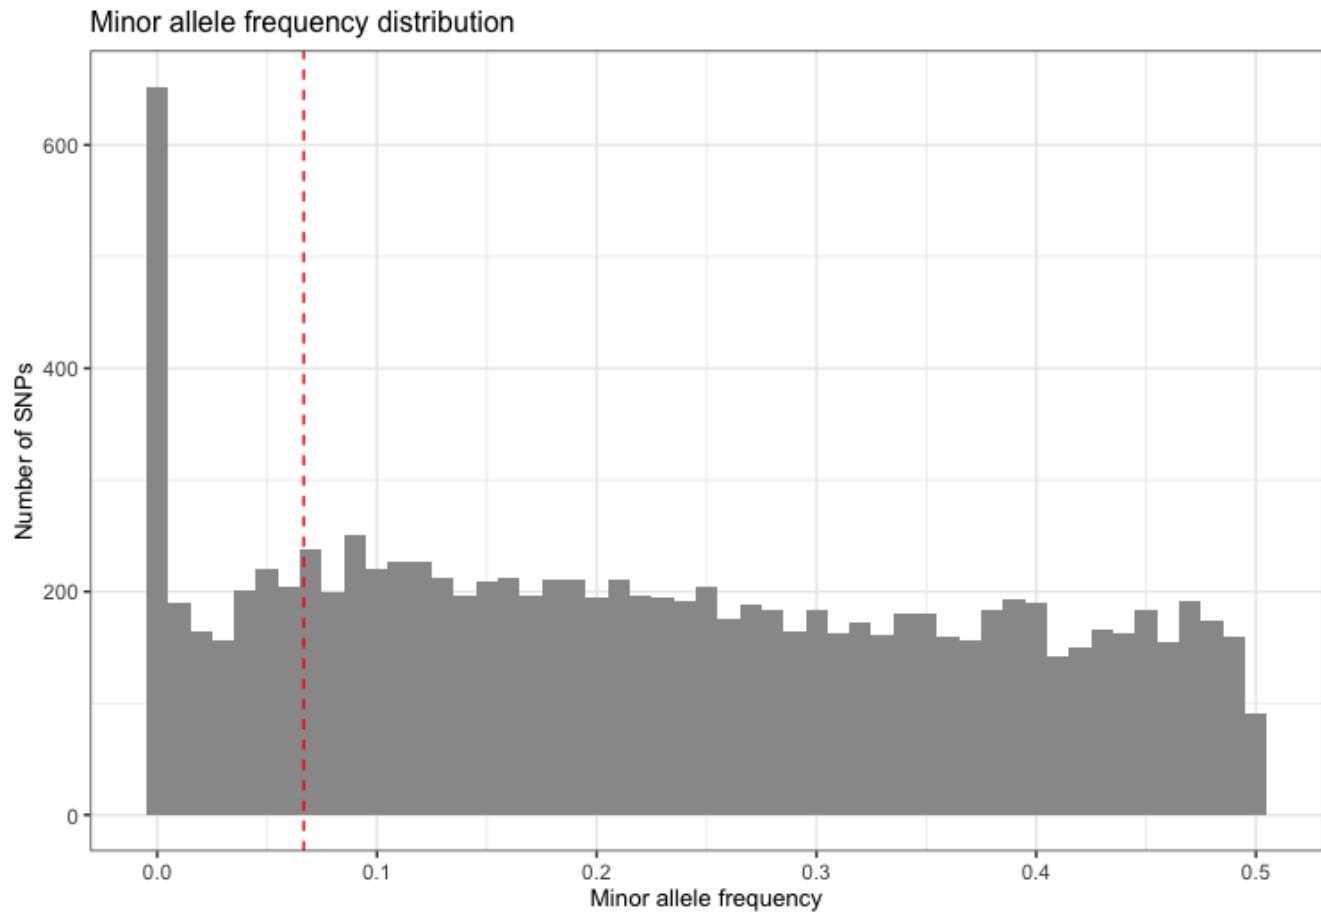
Distribution of $-\log_{10}(p\text{-value}_{\text{HWE}})$



Markers with low minor allele frequency

Markers with low minor allele count are often removed as the actual genotype calling (via the calling algorithm) is very difficult due to the small sizes of the heterozygote and rare-homozygote clusters. Identifying markers with low minor allele count is implemented in `check_maf`. It calculates the minor allele frequencies for all variants in the individuals that passed the `perIndividualQC`.

```
fail_maf <- check_maf(indir=indir, qcdir=qcdir, name=name, interactive=TRUE,
path2plink=path2plink, showPlinkOutput=FALSE)
```



References

1. Anderson CA, Pettersson FH, Clarke GM, Cardon LR, Morris AP, Zondervan KT. Data quality control in genetic case-control association studies. *Nature Protocols.* 2010;5: 1564–73. doi:10.1038/nprot.2010.116
2. Purcell S, Neale B, Todd-Brown K, Thomas L, Ferreira MAR, Bender D, et al. PLINK: a tool set for whole-genome association and population-based linkage analyses. *The American Journal of Human Genetics.* 2007;81: 559–75. doi:10.1086/519795
3. Chang CC, Chow CC, Tellier LC, Vattikuti S, Purcell SM, Lee JJ. Second-generation PLINK: rising to the challenge of larger and richer datasets. *GigaScience.* 2015;4: 7. doi:10.1186/s13742-015-0047-8
4. PLINK/SEQ. 2014. Available: <https://atgu.mgh.harvard.edu/plinkseq/index.shtml>
5. The International HapMap Consortium. A haplotype map of the human genome. *Nature.* 2005;437: 1299–320. doi:10.1038/nature04226
6. The International HapMap Consortium. A second generation human haplotype map of over 3.1 million SNPs. *Nature.* 2007;449: 851–. doi:10.1038/nature06258
7. The International HapMap Consortium. Integrating common and rare genetic variation in diverse human populations. *Nature.* 2010;467. doi:10.1038/nature09298
8. 1000 Genomes Project Consortium. An integrated map of structural variation in 2,504 human genomes. *Nature.* 2015;526: 75–81. doi:10.1038/nature15394

9. 1000 Genomes Project Consortium. A global reference for human genetic variation. *Nature*. 2015;526:75–81. doi:10.1038/nature15393